

## ✓ Assignment 4

### Text and sequence

```
from keras.layers import Embedding
# The Embedding layer takes at least two arguments:
# the number of possible tokens, here 1000 (1 + maximum word index),
# and the dimensionality of the embeddings, here 64.
embedding_layer = Embedding(1000, 64)

from keras.models import Sequential
from keras.layers import Flatten, Dense
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.callbacks import ModelCheckpoint
from keras.models import Sequential
from keras.layers import Flatten, Dense, Embedding, LSTM, Conv1D, MaxPooling1D,
from keras.models import load_model
from keras.preprocessing.text import Tokenizer
from sklearn.model_selection import train_test_split
from keras.optimizers import RMSprop
from google.colab import files
import re, os
from keras.datasets import imdb
from keras import preprocessing
from keras.utils import pad_sequences
```

### Model 1 From Scratch

```
# Number of words to consider as features
max_features = 10000
# Cut texts after 150 words
maxlen = 150
# Load the data as lists of integers.
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
#preprocessing.sequence.pad_sequences
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-data/17464789/17464789> [=====] - 1s 0us/step

```
model = Sequential()

# We specify the maximum input length to our Embedding layer
# so we can later flatten the embedded inputs

model.add(Embedding(10000, 8, input_length=maxlen))

# We flatten the 3D tensor of embeddings
# into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())

# We add the classifier on top
model.add(Dense(1, activation='sigmoid'))

#compiling the model
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history_1 = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_split=0.2)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 150, 8)	80000
flatten (Flatten)	(None, 1200)	0
dense (Dense)	(None, 1)	1201

Total params: 81201 (317.19 KB)  
Trainable params: 81201 (317.19 KB)  
Non-trainable params: 0 (0.00 Byte)

```
Epoch 1/10
625/625 [=====] - 26s 39ms/step - loss: 0.5974 - acc: 0.0000
Epoch 2/10
625/625 [=====] - 6s 9ms/step - loss: 0.3316 - acc: 0.0000
Epoch 3/10
625/625 [=====] - 4s 7ms/step - loss: 0.2566 - acc: 0.0000
Epoch 4/10
625/625 [=====] - 4s 6ms/step - loss: 0.2211 - acc: 0.0000
Epoch 5/10
625/625 [=====] - 4s 6ms/step - loss: 0.1968 - acc: 0.0000
Epoch 6/10
625/625 [=====] - 3s 5ms/step - loss: 0.1769 - acc: 0.0000
Epoch 7/10
625/625 [=====] - 3s 4ms/step - loss: 0.1595 - acc: 0.0000
Epoch 8/10
625/625 [=====] - 2s 4ms/step - loss: 0.1428 - acc: 0.0000
Epoch 9/10
625/625 [=====] - 2s 4ms/step - loss: 0.1267 - acc: 0.0000
```

Epoch 10/10

625/625 [=====] - 3s 5ms/step - loss: 0.1112 - acc: (

```
import matplotlib.pyplot as plt

accuracy = history_1.history['acc']
val_accuracy = history_1.history['val_acc']
loss = history_1.history['loss']
val_loss = history_1.history['val_loss']



epochs = range(1, len(accuracy) + 1)

plt.plot(epochs, accuracy, 'grey', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```



```
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.3465 - acc: 0.8657
Test loss: 0.34654125571250916
Test accuracy: 0.8657600283622742
```

Model 2 Training - 100 samples

```
max_features=10000
maxlen=150
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

texts = np.concatenate((x_train, x_test), axis=0)
labels = np.concatenate((x_train, x_test), axis=0)

x_train = x_train[:100]
y_train = y_train[:100]

model = Sequential()
model.add(Embedding(10000, 8, input_length=maxlen))
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
history_2 = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_split=0.2)
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 150, 8)	80000
flatten_1 (Flatten)	(None, 1200)	0
dense_1 (Dense)	(None, 1)	1201

=====  
Total params: 81201 (317.19 KB)  
Trainable params: 81201 (317.19 KB)  
Non-trainable params: 0 (0.00 Byte)

Epoch 1/10  
3/3 [=====] - 2s 440ms/step - loss: 0.6980 - acc: 0.4  
Epoch 2/10  
3/3 [=====] - 1s 243ms/step - loss: 0.6734 - acc: 0.4  
Epoch 3/10  
3/3 [=====] - 1s 276ms/step - loss: 0.6554 - acc: 0.4  
Epoch 4/10  
3/3 [=====] - 1s 286ms/step - loss: 0.6392 - acc: 0.4  
Epoch 5/10  
3/3 [=====] - 1s 226ms/step - loss: 0.6239 - acc: 1.0  
Epoch 6/10  
3/3 [=====] - 0s 175ms/step - loss: 0.6083 - acc: 1.0  
Epoch 7/10  
3/3 [=====] - 1s 207ms/step - loss: 0.5924 - acc: 1.0

```
Epoch 8/10
3/3 [=====] - 0s 171ms/step - loss: 0.5767 - acc: 1.0
Epoch 9/10
3/3 [=====] - 1s 189ms/step - loss: 0.5602 - acc: 1.0
Epoch 10/10
3/3 [=====] - 0s 142ms/step - loss: 0.5435 - acc: 1.0
```

```
accuracy = history_2.history['acc']
val_accuracy = history_2.history['val_acc']
loss = history_2.history['loss']
val_loss = history_2.history['val_loss']

epochs = range(1, len(accuracy) + 1)

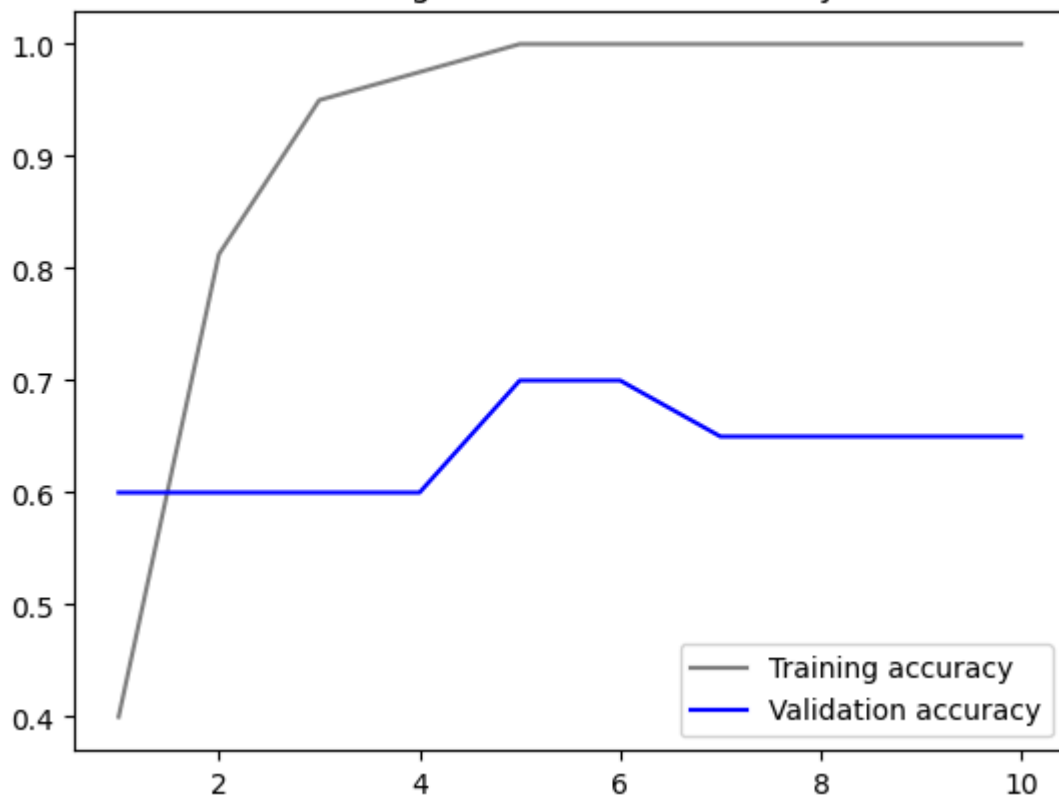
plt.plot(epochs, accuracy, 'grey', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

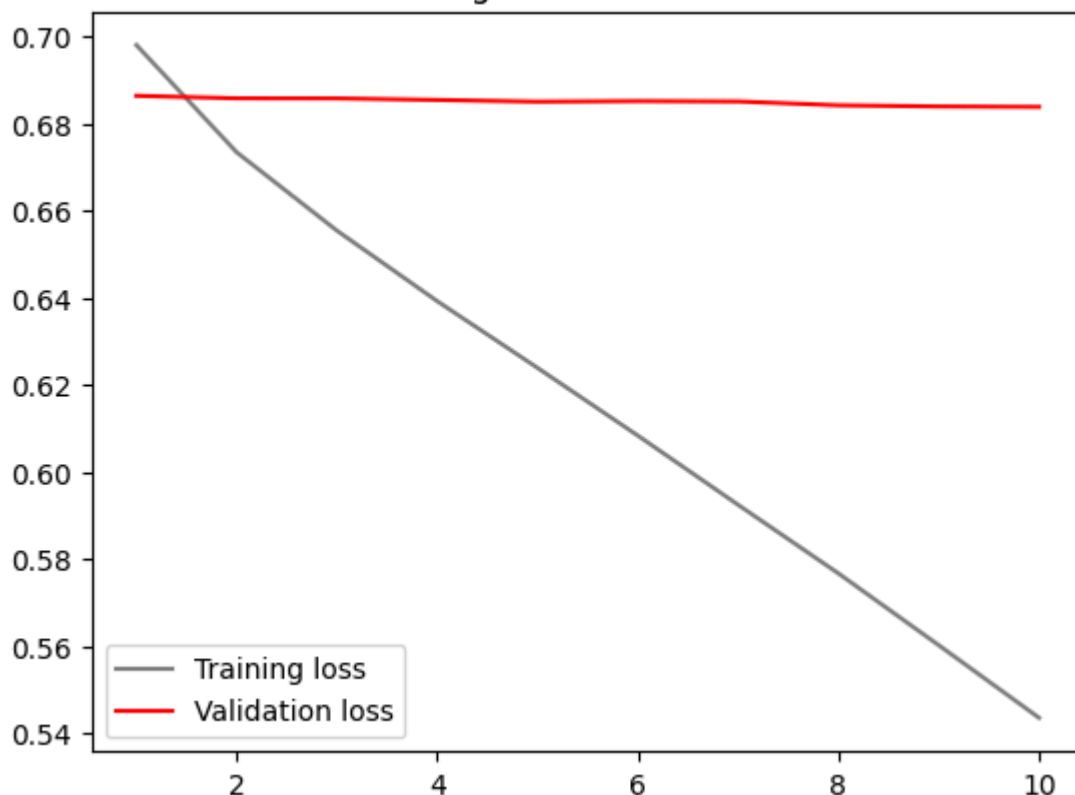
plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```

Training and validation accuracy



Training and validation loss



```
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

```
782/782 [=====] - 2s 3ms/step - loss: 0.6944 - acc: 0.4937
Test loss: 0.6943621039390564
Test accuracy: 0.49375998973846436
```

## ✓ Using Pre-Trained word embeddings

### Download the IMDB data as raw text

### Model 3 Pre-Trained model, Training- 100 samples

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

%cd /content/drive/MyDrive/

/content/drive/MyDrive

import os

!curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -xvf aclImdb_v1.tar.gz

!rm -r aclImdb/train/unsup
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Curre
			Dload Upload	Total	Spent	Left	Speed
100 80.2M	100 80.2M	0 0	7175k 0	0:00:11	0:00:11	--:--:--	14.3M

```
imdb_dir = '/content/drive/MyDrive/aclImdb'

train_dir = os.path.join(imdb_dir, 'train')

labels = []
texts = []

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(train_dir, label_type)
    for fname in os.listdir(dir_name):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)
```



## Tokenizing the data

```

maxlen = 150 # We will cut reviews after 100 words
training_samples = 100 # We will be training on 100 samples
validation_samples = 10000 # We will be validating on 10000 samples
max_words = 10000 # We will only consider the top 10,000 words in the dataset

tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))

data = pad_sequences(sequences, maxlen=maxlen)

labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)

# Split the data into a training set and a validation set
# But first, shuffle the data, since we started from data
# where sample are ordered (all negative first, then all positive).
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]

x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]

    Found 88582 unique tokens.
    Shape of data tensor: (25000, 150)
    Shape of label tensor: (25000,)

```

## Download the GloVe word embeddings

### Pre-Processing the embeddings

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call

```

from google.colab import drive
import os
import numpy as np

# Mount Google Drive
drive.mount('/content/drive')

# Specify the path to the GloVe embeddings file in your Google Drive
glove_dir = '/content/drive/MyDrive/glove.6B'
drive.mount("/content/drive", force_remount=True)
glove_file_path = os.path.join(glove_dir, 'glove.6B.100d.txt')

# Load GloVe embeddings into a dictionary
embeddings_index = {}
with open(glove_file_path, 'r', encoding='utf-8') as f:
    for line in f:
        values = line.split()
        word = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs

print('Found %s word vectors.' % len(embeddings_index))

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call  
Mounted at /content/drive

```

embedding_dim = 100

embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if i < max_words:
        if embedding_vector is not None:
            # Words not found in embedding index will be all-zeros.
            embedding_matrix[i] = embedding_vector

```

## Building the model

```

from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
model.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		

embedding_3 (Embedding)	(None, 150, 100)	1000000
lstm (LSTM)	(None, 32)	17024
dense_2 (Dense)	(None, 1)	33

```
=====
Total params: 1017057 (3.88 MB)
Trainable params: 1017057 (3.88 MB)
Non-trainable params: 0 (0.00 Byte)
```

---

## Loading the GloVe embeddings in the model

```
model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False
```

```
print("Training data shape:", y_train.shape)
```

```
Training data shape: (100,)
```

## Train and evaluate

```
import tensorflow as tf

# Check if Google Drive is mounted
if not tf.io.gfile.exists('/content/drive'):
    from google.colab import drive
    drive.mount('/content/drive')

# Retry saving the model weights
try:
    model.save_weights('/content/drive/MyDrive/pre_trained_glove_model.3a')
    print("Model weights saved successfully.")
except Exception as e:
    print("Error occurred while saving model weights:", e)
```

```
Model weights saved successfully.
```

```
import matplotlib.pyplot as plt

acc = history_3.history['acc']
val_acc = history_3.history['val_acc']
loss = history_3.history['loss']
val_loss = history_3.history['val_loss']

epochs = range(1, len(acc) + 1)


plt.plot(epochs, acc, 'grey', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()


plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```

Training and validation accuracy



Training and validation loss



```

test_dir = os.path.join(imdb_dir, 'test')

labels = []
texts = []

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)

sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)

try:
    model.load_weights('pre_trained_glove_model.3a')
    print("Model weights loaded successfully.")
except Exception as e:
    print("Error occurred while loading model weights:", e)

```

Error occurred while loading model weights: Unsuccessful TensorSliceReader co

## Now we change the number of training samples to

- ✓ **determine at what point the embedding layer gives better performance**

**Model 4 training sample size - 1000 using embedding layer**

```

max_features=10000
maxlen=150
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

texts = np.concatenate((x_train, x_test), axis=0)
labels = np.concatenate((x_train, x_test), axis=0)

x_train = x_train[:1000]
y_train = y_train[:1000]

model = Sequential()
model.add(Embedding(10000, 8, input_length=maxlen))
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
history_4 = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_split=0.2)

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 150, 8)	80000
flatten_2 (Flatten)	(None, 1200)	0
dense_3 (Dense)	(None, 1)	1201

=====  
Total params: 81201 (317.19 KB)

Trainable params: 81201 (317.19 KB)

Non-trainable params: 0 (0.00 Byte)

Epoch 1/10

25/25 [=====] - 2s 62ms/step - loss: 0.6934 - acc: 0

Epoch 2/10

25/25 [=====] - 1s 51ms/step - loss: 0.6765 - acc: 0

Epoch 3/10

25/25 [=====] - 1s 39ms/step - loss: 0.6595 - acc: 0

Epoch 4/10

25/25 [=====] - 1s 45ms/step - loss: 0.6382 - acc: 0

Epoch 5/10

25/25 [=====] - 2s 62ms/step - loss: 0.6118 - acc: 0

Epoch 6/10

25/25 [=====] - 1s 48ms/step - loss: 0.5799 - acc: 0

Epoch 7/10

25/25 [=====] - 1s 26ms/step - loss: 0.5432 - acc: 0

Epoch 8/10

25/25 [=====] - 1s 29ms/step - loss: 0.5025 - acc: 0

Epoch 9/10

25/25 [=====] - 1s 25ms/step - loss: 0.4591 - acc: 0

Epoch 10/10

25/25 [=====] - 0s 17ms/step - loss: 0.4141 - acc: 0

```
accuracy = history_4.history['acc']
val_accuracy = history_4.history['val_acc']
loss = history_4.history['loss']
val_loss = history_4.history['val_loss']

epochs = range(1, len(accuracy) + 1)



plt.plot(epochs, accuracy, 'grey', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```





```
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.6711 - acc: 0.5906
Test loss: 0.6710591912269592
Test accuracy: 0.5905600190162659
```

Double-click (or enter) to edit

## Model 5 Taining sample - 15000 using both embedding layer and Conv1D

```
max_features=10000
maxlen=150
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

texts = np.concatenate((x_train, x_test), axis=0)
labels = np.concatenate((x_train, x_test), axis=0)

x_train = x_train[:15000]
y_train = y_train[:15000]

model = Sequential()
model.add(Embedding(10000, 10, input_length=maxlen))
model.add(Conv1D(512, 3, activation='relu'))
model.add(MaxPooling1D(3))

model.add(Conv1D(256, 3, activation='relu'))
model.add(MaxPooling1D(3))

model.add(Conv1D(256, 3, activation='relu'))
model.add(Dropout(0.8))
model.add(MaxPooling1D(3))

model.add(GlobalMaxPooling1D())
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
history_5 = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_split=0.2)
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
embedding_5 (Embedding)	(None, 150, 10)	100000
conv1d (Conv1D)	(None, 148, 512)	15872
max_pooling1d (MaxPooling1D)	(None, 49, 512)	0
conv1d_1 (Conv1D)	(None, 47, 256)	393472

max_pooling1d_1 (MaxPooling1D)	(None, 15, 256)	0
conv1d_2 (Conv1D)	(None, 13, 256)	196864
dropout (Dropout)	(None, 13, 256)	0
max_pooling1d_2 (MaxPooling1D)	(None, 4, 256)	0
global_max_pooling1d (GlobalMaxPooling1D)	(None, 256)	0
flatten_3 (Flatten)	(None, 256)	0
dense_4 (Dense)	(None, 1)	257

```

=====
Total params: 706465 (2.69 MB)
Trainable params: 706465 (2.69 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

Epoch 1/10
375/375 [=====] - 27s 56ms/step - loss: 0.6919 - acc: 0.0000
Epoch 2/10
375/375 [=====] - 8s 21ms/step - loss: 0.4760 - acc: 0.0000
Epoch 3/10
375/375 [=====] - 5s 12ms/step - loss: 0.3332 - acc: 0.0000
Epoch 4/10
375/375 [=====] - 5s 14ms/step - loss: 0.2775 - acc: 0.0000
Epoch 5/10
375/375 [=====] - 4s 10ms/step - loss: 0.2394 - acc: 0.0000
Epoch 6/10
375/375 [=====] - 3s 8ms/step - loss: 0.2065 - acc: 0.0000
Epoch 7/10
375/375 [=====] - 5s 12ms/step - loss: 0.1832 - acc: 0.0000
Epoch 8/10
375/375 [=====] - 4s 11ms/step - loss: 0.1537 - acc: 0.0000
Epoch 9/10
375/375 [=====] - 3s 8ms/step - loss: 0.1328 - acc: 0.0000
Epoch 10/10
375/375 [=====] - 3s 7ms/step - loss: 0.1128 - acc: 0.0000

```

```
accuracy = history_5.history['acc']
val_accuracy = history_5.history['val_acc']
loss = history_5.history['loss']
val_loss = history_5.history['val_loss']

epochs = range(1, len(accuracy) + 1)


plt.plot(epochs, accuracy, 'grey', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()


plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```

Training and validation accuracy



Training and validation loss



```
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

```
782/782 [=====] - 3s 4ms/step - loss: 0.4450 - acc: 0.7938
Test loss: 0.44504714012145996
Test accuracy: 0.7938799858093262
```

As we have seen in the previous model even though we increased the training sample size the accuracy was still low but when we used Conv1D along with increased training sample size the accuracy improved to 81%

### Model 6 Training sample 30000 using both embedding layers and Conv1D

```
max_features=10000
maxlen=150
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)

x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

texts = np.concatenate((x_train, x_test), axis=0)
labels = np.concatenate((x_train, x_test), axis=0)

x_train = x_train[:30000]
y_train = y_train[:30000]

model = Sequential()
model.add(Embedding(10000, 12, input_length=maxlen))
model.add(Conv1D(512, 3, activation='relu'))
model.add(MaxPooling1D(3))

model.add(Conv1D(256, 3, activation='relu'))
model.add(MaxPooling1D(3))

model.add(Conv1D(256, 3, activation='relu'))
model.add(Dropout(0.8))
model.add(MaxPooling1D(3))

model.add(GlobalMaxPooling1D())
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
history_6 = model.fit(x_train, y_train,
                      epochs=10,
                      batch_size=32,
                      validation_split=0.2)
```

Model: "sequential\_10"

Layer (type)	Output Shape	Param #
embedding_11 (Embedding)	(None, 150, 12)	120000
conv1d_12 (Conv1D)	(None, 148, 512)	18944
max_pooling1d_12 (MaxPooling1D)	(None, 49, 512)	0

conv1d_13 (Conv1D)	(None, 47, 256)	393472
max_pooling1d_13 (MaxPooling1D)	(None, 15, 256)	0
conv1d_14 (Conv1D)	(None, 13, 256)	196864
dropout_4 (Dropout)	(None, 13, 256)	0
max_pooling1d_14 (MaxPooling1D)	(None, 4, 256)	0
global_max_pooling1d_4 (GlobalMaxPooling1D)	(None, 256)	0
flatten_7 (Flatten)	(None, 256)	0
dense_10 (Dense)	(None, 1)	257

```

=====
Total params: 729537 (2.78 MB)
Trainable params: 729537 (2.78 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

Epoch 1/10
625/625 [=====] - 18s 28ms/step - loss: 0.6481 - acc: 0.0000
Epoch 2/10
625/625 [=====] - 8s 12ms/step - loss: 0.3968 - acc: 0.0000
Epoch 3/10
625/625 [=====] - 5s 9ms/step - loss: 0.3205 - acc: 0.0000
Epoch 4/10
625/625 [=====] - 6s 9ms/step - loss: 0.2810 - acc: 0.0000
Epoch 5/10
625/625 [=====] - 4s 7ms/step - loss: 0.2525 - acc: 0.0000
Epoch 6/10
625/625 [=====] - 4s 7ms/step - loss: 0.2262 - acc: 0.0000
Epoch 7/10
625/625 [=====] - 5s 9ms/step - loss: 0.1999 - acc: 0.0000
Epoch 8/10
625/625 [=====] - 5s 7ms/step - loss: 0.1770 - acc: 0.0000
Epoch 9/10
625/625 [=====] - 4s 7ms/step - loss: 0.1540 - acc: 0.0000
Epoch 10/10
625/625 [=====] - 5s 8ms/step - loss: 0.1278 - acc: 0.0000

```

```
accuracy = history_6.history['acc']
val_accuracy = history_6.history['val_acc']
loss = history_6.history['loss']
val_loss = history_6.history['val_loss']

epochs = range(1, len(accuracy) + 1)



plt.plot(epochs, accuracy, 'grey', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'grey', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```





```
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test loss:', test_loss)
print('Test accuracy:', test_acc)
```

```
782/782 [=====] - 3s 4ms/step - loss: 0.4130 - acc: 0.8128
Test loss: 0.41296547651290894
Test accuracy: 0.8127999901771545
```

## Model 7 pretrained model. Training - 15000 samples

```
maxlen = 150 # We will cut reviews after 150 words
training_samples = 15000 # We will be training on 15000 samples
validation_samples = 10000 # We will be validating on 10000 samples
max_words = 10000 # We will only consider the top 10,000 words in the dataset

tokenizer = Tokenizer(num_words=max_words)

# Convert any non-string elements in the texts list to strings
texts = [str(text) for text in texts]

# Fit the tokenizer on the texts
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))

data = pad_sequences(sequences, maxlen=maxlen)

labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)

# Split the data into a training set and a validation set
# But first, shuffle the data, since we started from data
# where samples are ordered (all negative first, then all positive).
indices = np.arange(data.shape[0])
np.random.shuffle(indices)

data = data[indices]
```